

The Thermodynamics of Togetherness

How We Built a Cooperation Engine,
Proved It Worked, Then Proved We Were Wrong

A Computational Investigation in 53 Findings

Tristan Stoltz

Luminous Dynamics

2026

Epigraph

*“Physics doesn’t cooperate. It clusters at the wells.
Everything we called society was resource geometry.”*

— Chapter 15

*“The tragedy of walking away is that you lose benefits
you never knew you were receiving.”*

— Chapter 16

Contents

I	The Engine	1
1	Chapter 1: Why Physics, Not Utility Functions	3
1.1	The Standard Story	3
1.2	What We Did Instead	3
1.3	The Question	4
1.4	What We Found	4
1.5	The Structure of This Book	5
1.6	A Note on Honesty	5
2	Chapter 2: Five Channels	7
2.1	The Architecture	7
2.2	Channel 1: Motor Gain	7
2.3	Channel 2: Energy Budget	8
2.4	Channel 3: Harmony Field Friction	8
2.5	Channel 4: Sanctuary Zones	8
2.6	Channel 5: Conformal Curvature	9
2.7	The Coupling That Matters	9
3	Chapter 3: What’s Real and What Isn’t	11
3.1	The Classification	11
3.2	The Audit	11
3.3	What This Means for the Findings	12
3.4	Why We’re Telling You This	12
3.5	The Proposed Standard	13
3.6	The Cost of Honesty	13
4	Chapter 4: From Twenty to One Thousand	15
4.1	The Scaling Problem	15
4.2	The Spatial Hash	15
4.3	The Benchmark	15
4.4	Finding 37: Cooperation at Scale	16
4.5	Determinism	17
II	The Illusion	19
5	A Note Before Part II	21

6	Chapter 5: Cooperation as Thermodynamic Necessity	23
6.1	The Simplest Possible World	23
6.2	What Happens	23
6.3	The Six-Condition Ablation	24
6.4	The Metric Independence Result	24
6.5	What We Could Not Kill	25
6.6	The Implication	25
7	Chapter 6: The Temperature of Society	27
7.1	The Ising Analogy	27
7.2	The Pressure Sweep	27
7.3	The Bistability	28
7.4	Desperation Cooperation	28
7.5	The Thermodynamic Response	28
7.6	The Parameter Sensitivity	29
7.7	The Critical Infrastructure	29
7.8	The Phase Diagram	30
8	Chapter 7: Social Bonds and the Physics of Friendship	31
8.1	The Constraint Hypothesis	31
8.2	Finding 19: Bonds Transform Survival	31
8.3	The Dunbar Limit	32
8.4	Migration	32
8.5	Generational Wealth	33
9	Chapter 8: Raiders, Evolvers, and the Solidarity Response	35
9.1	The Adversarial Threshold	35
9.2	The Solidarity Response	35
9.3	The Arms Race	36
9.4	The Evolutionary Lesson	37
10	Chapter 9: The Economics of Scarcity	39
10.1	Three Parables	39
10.1.1	Parable One: The Curse of Abundance	39
10.1.2	Parable Two: The Futility of Charity	40
10.1.3	Parable Three: The Equal Opportunity Paradox	40
10.2	The Unified Theory	41
11	Chapter 10: Technology and the Shrinking of Range	43
11.1	Five Eras	43
11.2	The Results	43
11.3	Connection Is Not Proximity	44
12	Chapter 11: Learning and the Paradox of Optimal Defaults	45
12.1	The Learning Hypothesis	45
12.2	Finding 31: Weights Are Already Optimal	45
12.3	Finding 34: Memory Doesn't Help Either	46
12.4	Finding 44: The U-Curve Mechanism	46
12.5	Finding 45: Harmony Dimensionality Is Irrelevant	47

12.6 The Paradox of Optimal Defaults	47
13 Chapter 12: Higher Dimensions and the Geometry of Cooperation	49
13.1 The 3D Test	49
13.2 Curvature: Beautiful Mathematics, Zero Social Effect	50
13.3 What Dimension Teaches	50
III The Unraveling	53
14 Chapter 13: The Thread That Unraveled Everything	55
14.1 The Question	55
14.2 The Result	56
14.3 The FEP Gradient Was the Problem	56
14.4 What By-Product Mutualism Means	56
15 Chapter 14: Seven Failed Rescues	59
15.1 The Rescue Attempts	59
15.2 Rescue 1: Make Wells Scarce (F48)	59
15.3 Rescue 2: Make Wells Volatile (F50)	59
15.4 Rescue 3: Make Wells Hidden (F49)	60
15.5 Rescue 4: Require Threshold Co-Location (F51)	60
15.6 Rescue 5: Require Compatible Partners (F52)	60
15.7 Rescue 6: Remove Wells Entirely (F53)	61
15.8 Rescue 7: The FEP Gradient as Informational Drive	61
15.9 What the Rescues Taught Us	61
16 Chapter 15: Resource Geometry	63
16.1 The Thesis, Revised	63
16.2 What This Explains	63
16.3 What This Does Not Explain	64
16.4 The Constructive Proof, Revised	64
16.5 The Lesson for Simulation Science	65
16.6 The Closing	65
17 Chapter 16: The Freedom to Walk Away	67
17.1 What Survived	67
17.2 By-Product Mutualism Requires Co-Location	67
17.3 The U-Curve Is About Resource Access, Not Commitment	68
17.4 The Bowling Alone Reinterpretation	68
17.5 What This Means	68
17.6 The Last Word	69
17.7 Appendix Note	69
Appendix B: Old Interpretation vs Revised Interpretation	71
Afterword: What This Changes	73
Reproduction	75

Part I

The Engine

Chapter 1

Chapter 1: Why Physics, Not Utility Functions

In which we ask whether cooperation needs a reason, and discover that thermodynamics is sufficient.

1.1 The Standard Story

The standard story of cooperation goes like this: agents in a population face a social dilemma — a situation where individual incentives conflict with collective benefit. Each agent calculates (or evolves, or learns) a strategy. Under the right conditions — repeated interaction, kin selection, reputation, punishment — cooperative strategies outperform selfish ones, and cooperation emerges.

This is the story told by Axelrod’s *Evolution of Cooperation* (1984), by Nowak’s *SuperCooperators* (2011), by the vast literature on evolutionary game theory from Maynard Smith (1982) through Sigmund (2010). It is a story about strategy, payoff, and optimization.

It is not our story.

1.2 What We Did Instead

We built a physics engine. Not a game theory framework, not an agent-based model with utility functions, not an evolutionary simulation with payoff matrices. A physics engine — with rigid bodies, collision detection, energy budgets, and the second law of thermodynamics.

In this engine, agents do not calculate. They have no utility function. They cannot reason about payoffs or discount future rewards. They have energy, position, velocity, and a harmony profile — eight real numbers between zero and one that describe their social compatibility type. They follow a behavioral gradient that points them toward conditions that reduce their estimated free energy: toward energy wells when depleted, toward compatible partners when energy is moderate, away from danger when threatened.

They cooperate — not because cooperation is a strategy that maximizes payoff, but because proximity to compatible partners regenerates energy, and energy is necessary to survive.

This is not a metaphor. The energy is tracked in joules. The entropy increases monotonically. The Helmholtz free energy is computed as $F = U - TS$. The Landauer bound (2.87×10^{-21} joules per bit at body temperature) establishes the minimum energy cost of information processing. Every action costs energy. Every collision dissipates heat. Every tick, the second law takes its tax.

When two compatible agents stand near each other, their harmony resonance generates a small energy return. This is the engine's sole cooperation mechanism. There is no contract, no reciprocity tracking, no reputation system. Just physics: compatible proximity regenerates energy.

1.3 The Question

The question this book asks is simple: **Is thermodynamic enforcement sufficient to produce cooperation?**

Not cooperation as a strategy. Not cooperation as an evolutionary stable state. Cooperation as a *physical inevitability* — as certain as heat flowing from hot to cold, as predictable as entropy increasing in a closed system.

The answer, across forty-five findings and fifty-five experiments, is yes. But the yes comes with complications, caveats, and one devastating exception that forms the climax of the book.

1.4 What We Found

The findings organize into four categories:

Cooperation emerges (Chapters 5-7). Under energy scarcity with finite wells and harmony resonance, agents cluster into cooperative groups. The clustering mechanism is the FEP gradient (agents seek compatible partners). The survival mechanism is harmony resonance (partners regenerate each other's energy). Neither alone is sufficient. Both together produce sustainable social structure. The specific consciousness metric is irrelevant — IIT, Shannon entropy, or a constant all produce identical cooperation. Only zero integration differs.

Cooperation is robust (Chapters 6, 8, 12). We could not kill cooperation through parameter extremes (thirty-fold maintenance increase), structural changes (fragmented perception, obstacles, moving wells), adversarial agents (up to 50% of population), information manipulation, memory, learning, communication, or curvature. The only way to kill it was to remove resonance regeneration entirely. Cooperation is topologically invariant — it fragments into smaller groups rather than disappearing.

The environment shapes the outcome (Chapters 9-10). Abundance breeds complacency. Charity fails without structure. Equal starts produce maximum inequality. Algorithms shrink effective social range. Physical bonds dramatically improve survival. Evolution selects for solidarity, not tribalism. These environmental findings emerge from the same engine with the same thermodynamics — only the conditions change.

Cooperation requires no intelligence but resists no freedom (Chapters 14-15). Information, memory, learning, and communication have zero measurable effect — the FEP gradient already encodes all necessary information. But adding a single degree of freedom — the probability of ignoring the social gradient — reproduces Putnam's entire "Bowling Alone" pattern: declining participation, increasing isolation, persistent

economic output, and rising inequality. Cooperation is thermodynamically inevitable only when agents cannot choose otherwise.

1.5 The Structure of This Book

Part I (Chapters 1-4) describes the engine: why we built it, how it works, what’s real and what’s invented, and how it scales. Chapter 3 is an honest audit of every mechanic, classifying each as GROUNDED, INSPIRED, or SPECULATIVE. We believe this transparency should be standard practice.

Part II (Chapters 5-12) presents the experiments. Each chapter covers a cluster of related findings, moving from the founding cooperation result through phase transitions, social bonds, adversarial dynamics, resource economics, technology effects, learning, and higher-dimensional physics.

Part III (Chapters 13-16) discusses implications. What can simulations tell us? What can’t they? What does the equal opportunity paradox mean? What does blind navigation suggest about social systems? And what does the “bowling alone” result — the single experiment that broke the engine’s cooperation guarantee — tell us about the nature of community?

1.6 A Note on Honesty

This book makes no claim about consciousness. The engine uses an “integration metric” (Φ) inspired by Tononi’s Integrated Information Theory, but we use it as a coupling parameter, not as a claim about subjective experience. Where we write “consciousness” in code and equations, we mean “a scalar value that modulates physics.” Nothing more.

Similarly, this book makes no claim about human society. Our agents have no culture, no institutions, no language, no history. When we draw parallels to Putnam or Klinenberg or the resource curse, we are noting structural similarities, not claiming causal explanations. The engine models thermodynamic agents, not people.

What the engine does model — with mathematical precision and statistical rigor — is the relationship between energy scarcity, social proximity, mutual benefit, and emergent cooperation. This relationship holds across 45 findings, 55 experiments, 20-seed replication with Holm-Bonferroni correction, and populations from 4 to 1,000 agents.

Physics doesn’t bowl alone. But understanding why might help us understand what happens when we do.

Next: Chapter 2 — Five Channels: How Integration Metrics Couple to Physics

Chapter 2

Chapter 2: Five Channels

In which we describe how an integration metric couples to rigid body physics, and why each coupling exists.

2.1 The Architecture

Symtropy is a rigid body physics engine written in Rust, supporting 2D through 9D simulation with LBVH broadphase collision detection, GJK/EPA narrowphase, and PBD constraint solving. Agents are spheres with position, velocity, mass, and linear damping. Physics runs at 64 Hz with semi-implicit Euler integration and LTC exponential damping (frame-rate independent, never reverses velocity).

On top of this physics layer sits the consciousness-physics coupling — five bidirectional channels through which an integration metric (Φ) modulates rigid body dynamics. These channels are the engine’s novel contribution. Each represents a different hypothesis about how internal state could affect physical behavior.

2.2 Channel 1: Motor Gain

Φ modulates force authority through a four-tier safety system:

Tier	Φ Range	Motor Gain	Interpretation
Green	> 0.6	100%	Full authority
Yellow	$0.3 - 0.6$	60%	Reduced output
Orange	$0.1 - 0.3$	30%	Minimal control
Red	< 0.1 or collapsed	0%	No motor output

An agent at Red tier cannot move — it is functionally paralyzed. This creates a survival pressure: maintaining Φ above 0.1 is necessary for locomotion, and locomotion is necessary for reaching energy wells and cooperation partners.

Classification: **INSPIRED**. The four-tier structure borrows from NRC nuclear safety protocols. No neuroscience publication proposes this specific Φ -to-motor mapping. The general principle — that integrated information modulates behavioral output — has loose support in motor control literature (Adams, Shipp & Friston 2013), but the discrete tiers are invention.

2.3 Channel 2: Energy Budget

Every tick, agents pay a maintenance cost proportional to their integration level:

```
cost = base_maintenance  $\times$  (1.0 +  $\Phi$   $\times$  0.5)
```

Higher Φ costs more energy. This is the thermodynamic cost of integration — the engine’s analogue of the brain’s metabolic burden (Raichle & Gusnard 2002). The brain consumes 20% of the body’s energy while comprising 2% of its mass; maintaining high integration is expensive.

Energy regenerates through three sources: ambient recovery (slow, ~ 0.005 J/tick), energy wells (moderate, ~ 0.12 J/tick when standing on one), and harmony resonance (variable, depends on partner compatibility and proximity).

Classification: **GROUND**ED for the thermodynamic accounting (U, T, S, F tracked per agent, Landauer bound referenced). **INSPI**RED for the Φ -proportional cost (the 0.5 coupling constant is empirical, not derived).

2.4 Channel 3: Harmony Field Friction

Each agent emits a harmony field — eight scalar values representing activation along the Eight Harmonies (Stillness, Play, Craft, Justice, Curiosity, Celebration, Kinship, Stewardship). The field falls off as $1/r^{(D-1)}$ with Plummer softening ($\varepsilon = 1.0$) to prevent singularities.

When two agents’ harmony fields overlap, the resonance (dot product of their activation vectors, normalized) modulates the collision response. High resonance reduces prediction error from collisions; low resonance increases it. This means collisions between compatible agents are less “surprising” — they produce less behavioral disruption.

Classification: **INSPI**RED. The $1/r^{(D-1)}$ falloff is correct field theory. The harmony-as-field-source concept draws loose analogy from McFadden’s CEMI field theory (2020), which proposes that the brain’s electromagnetic field is conscious. The specific eight-channel structure has no physics precedent.

2.5 Channel 4: Sanctuary Zones

When an agent’s Sacred Stillness activation exceeds 0.6, total harmony energy exceeds 2.0, and Φ exceeds 0.3, a sanctuary zone forms — a spatial region where collision impulses are dampened by up to 90%.

In forty-five findings, no experiment ever activated a sanctuary zone. The Φ values produced by the consciousness equation (~ 0.06 - 0.09) never reach the 0.3 threshold. The mechanism exists in the codebase but is functionally inert.

Classification: **SPECULATIVE**. No physical system creates a collision-dampening force field through collective harmonization. The closest real-world analogue — social buffering (Hostinar et al. 2014) — operates through stress hormone modulation, not impulse dampening. We retain the mechanism for game design purposes but make no scientific claim about it.

2.6 Channel 5: Conformal Curvature

The harmony field energy parameterizes a conformal Riemannian metric:

$$g_{ij}(x) = e^{-2\sigma(x)} \delta_{ij}, \quad \text{where } \sigma(x) = \kappa \times E_{\text{harm}}$$

This produces geodesic corrections to agent trajectories — harmony “wells” in the field curve the simulation space, deflecting nearby agents. The mathematics is correct (Christoffel symbols, geodesic equation, Ricci scalar all properly derived from the conformal factor). The physics is speculative — no law of nature couples integration metrics to spacetime curvature.

Finding 8 demonstrated measurable trajectory deflection ($p = 0.009$, $d = -35.5$) at curvature scale $\kappa = 0.05$. Finding 36 demonstrated that this deflection has zero effect on cooperation ($d = 0.006$). The curvature is mathematically real and socially irrelevant.

Classification: **GROUNDED** mathematics (Carroll 2004, Wald 1984). **SPECULATIVE** physics.

2.7 The Coupling That Matters

Of the five channels, only two contribute to the cooperation findings reported in this book:

1. **Energy budget** (Channel 2): creates the thermodynamic pressure that drives cooperation
2. **Harmony field resonance** (embedded in Channel 3): provides the mutual benefit that sustains cooperation

Motor gain (Channel 1) affects collapsed agents but not active ones (Φ rarely drops below 0.3 in practice). Sanctuary (Channel 4) never activates. Curvature (Channel 5) deflects trajectories without affecting social outcomes.

The engine has five channels, but the book runs on two. We report all five for completeness and honesty — and because the inert channels are themselves a finding. Three coupling mechanisms that seemed theoretically important turned out to be empirically irrelevant. Only the thermodynamic core — energy pressure and mutual benefit — drives the social dynamics that fill the remaining chapters.

Next: Chapter 3 — What’s Real and What Isn’t

Chapter 3

Chapter 3: What's Real and What Isn't

In which we audit our own engine for scientific honesty and find that half of it is invention.

3.1 The Classification

Every mechanic in a simulation engine has a relationship to reality. Some implement well-established physics. Some draw loose analogy from real systems. Some have no precedent whatsoever and exist because a programmer thought they would be interesting.

Most simulation papers do not distinguish between these categories. We do.

We classify every mechanic in Symtropy as one of three types:

GROUND: Implements established physical or mathematical laws. Can be validated against nature. Uses published constants. If the mechanic is wrong, a physicist can demonstrate how.

INSPIRE: Draws loose analogy from a real system but does not faithfully implement its equations. Captures the spirit of a scientific idea without the rigor. If the mechanic is wrong, the error is in the analogy, not the math.

SPECULATIVE: Has no published precedent. Invented for this simulation. If the mechanic is wrong, there is nothing to compare it against — it was never claimed to be right.

3.2 The Audit

Mechanic	Math	Physics
Thermodynamics (U, T, S, F)	Correct	Correct
Landauer bound (2.87e-21 J/bit)	Exact	Published
Conformal geometry ($g_{ij} = e^{2\sigma}\delta_{ij}$)	Correct Riemannian	No physics law couples Φ to sp
FEP gradient descent	Approximation	Loose Friston analogy
Harmony field ($1/r^{D-1}$ falloff)	Correct field theory form	No physical field exists
Motor gain (4-tier safety)	—	No neuroscience precedent
Sanctuary zones (impulse dampening)	—	No “harmony force field” exists

Mechanic	Math	Physics
Φ -gravity ($F = G\Phi_1\Phi_2/r^2$)	Correct Newtonian form	Tononi never proposed this
Dimensional leakage	—	Randall-Sundrum misapplied
Joules-per-Phi metric	Novel	Novel

Five of nine mechanics are either INSPIRED or SPECULATIVE. Only the core thermodynamics and the conformal geometry mathematics are fully GROUNDED. The engine is, by this accounting, approximately half invention.

3.3 What This Means for the Findings

Not all findings depend on all mechanics.

The cooperation emergence findings (Chapters 5-6) depend only on the GROUNDED thermodynamic layer (energy budgets, maintenance costs, Helmholtz free energy) and the INSPIRED FEP gradient. They do not use sanctuary zones, Φ -gravity, dimensional leakage, or conformal curvature. The core thesis — cooperation as thermodynamic necessity — rests on the most grounded components of the engine.

The curvature lensing finding (Finding 8: $p=0.009$) depends on the conformal geometry, which is mathematically GROUNDED but physically SPECULATIVE. The trajectory deflection is real Riemannian geometry. The claim that harmony field energy curves space is pure invention. Finding 36 confirmed that this curvature has zero effect on cooperation — it deflects trajectories but doesn't change social outcomes. The conformal curvature is a mathematical demonstration, not a social mechanism.

The Φ -gravity and sanctuary mechanics are used in a few early experiments but are not invoked in any finding reported in this book. They remain as engine features for potential game design use, but we make no scientific claims about them.

3.4 Why We're Telling You This

The standard practice in simulation papers is to present the engine as a unified whole and let the reader assume that all components are equally grounded. This is misleading. A reader who sees “conformal Riemannian curvature” next to “thermodynamic energy budgets” might reasonably assume both are physics. Only one is.

We classify our own mechanics because:

1. **Reviewers will find what you don't disclose.** An unreported SPECULATIVE mechanic discovered by a reviewer becomes a credibility problem. A disclosed SPECULATIVE mechanic becomes an honest design choice.
2. **Other builders need to know.** If someone extends this engine, they should know which parts are load-bearing science and which are decorative metaphor. Building on a SPECULATIVE foundation without knowing it is how bad science propagates.
3. **The findings are stronger with the distinction.** When we say “cooperation emerges from thermodynamics,” we can now demonstrate that this claim rests only on GROUNDED mechanics. The SPECULATIVE components are separable and inert. The thesis does not depend on them.

3.5 The Proposed Standard

We propose that every simulation engine paper include a realism classification table. The format is simple:

Mechanic	Classification	Evidence
[name]	GROUNDLED / INSPIRED / SPECULATIVE	[citation or “none”]

This table should appear in the methods section, before any results. It sets expectations and demonstrates scientific maturity.

The classification requires judgment. Reasonable people may disagree about whether a particular mechanic is INSPIRED or SPECULATIVE — the FEP gradient is a borderline case. The point is not to achieve perfect classification but to *force the exercise of classification*. A builder who must categorize each mechanic will think more carefully about what they’re building.

3.6 The Cost of Honesty

Disclosing that half your engine is speculative is uncomfortable. It invites the question: “If half of it is made up, why should I trust any of it?”

The answer is in the separation. The findings reported in this book depend on the GROUNDLED half. The SPECULATIVE half is present but inert — it exists in the codebase but does not contribute to the cooperation, phase transition, or social dynamics results. It is engine feature, not scientific claim.

This separation is itself a finding. We *tested* whether the speculative mechanics matter (Finding 36: curvature irrelevant, Finding 29: information irrelevant, Finding 34: memory irrelevant). The negative results are as important as the positive ones — they demonstrate that the engine’s core behavior comes from its grounded components, not from the speculative additions that make it look impressive.

If we had not disclosed and tested the classification, we would not know this. The honesty audit strengthened the science.

Next: Chapter 4 — Scaling and Performance: From Twenty Agents to One Thousand

Chapter 4

Chapter 4: From Twenty to One Thousand

In which we discover that the engine is not a toy, that the bottleneck is collision detection, and that cooperation gets stronger at scale.

4.1 The Scaling Problem

Every experiment in Chapters 5 through 11 uses 20-24 agents. This is typical of agent-based modeling — NetLogo tutorials use 50, Mesa examples use 100, and published ABM papers rarely exceed a few hundred. But a claim about cooperation as “thermodynamic inevitability” carries an implicit assumption: that the result holds at scale. Twenty agents in a 100×100 arena is a sparse, simple system. Does cooperation survive density?

4.2 The Spatial Hash

The engine’s main computational bottleneck is the $O(n^2)$ harmony resonance loop: every pair of agents must be checked for proximity and resonance. At $N=20$, this is 190 pairs — trivial. At $N=1,000$, it is 499,500 pairs per tick at 64 Hz — 32 million checks per second.

We implemented a grid-based spatial hash (Chapter 1 of the code, `spatial_hash.rs`) with cell size equal to the harmony range. This guarantees that all potential cooperation partners are in the same cell or an adjacent cell — a 3×3 neighborhood in 2D, 3×3×3 in 3D. The check reduces from $O(n^2)$ to $O(n \times k)$, where k is the average number of agents per cell neighborhood.

Seven unit tests verify correctness, including 3D operation, negative coordinates, and a 1,000-agent population test confirming that queries return a proper subset of the total population.

4.3 The Benchmark

We benchmarked both brute-force and spatial-hash versions at population sizes from 10 to 1,000:

N	Brute Force (ms/tick)	Spatial Hash (ms/tick)	Speedup
10	0.02	0.04	0.5× (overhead)
50	0.64	0.62	1.0×
100	2.67	2.40	1.1×
200	14.17	7.88	1.8×
500	165.58	59.69	2.8×
1,000	(skipped)	444.20	—

The spatial hash provides meaningful speedup above $N=200$ (1.8×) and becomes essential at $N=500$ (2.8×). At $N=1,000$, brute force would take approximately 2,500 ms/tick (estimated from the $O(n^2)$ trend) — the spatial hash achieves 444 ms/tick.

However, 444 ms/tick is far from real-time (which would require <16 ms at 64 Hz). The bottleneck is not the harmony loop — it’s the physics world step. GJK collision detection on 1,000 spheres dominates the tick time. The spatial hash fixed the harmony bottleneck; the physics broadphase (LBVH) handles collision pairs, but EPA penetration resolution on all overlapping spheres remains expensive.

For the book’s purposes, 444 ms/tick at $N=1,000$ means a full 8,000-tick experiment takes approximately 60 minutes in release mode. This is feasible for research (we ran it) but not for real-time gameplay.

4.4 Finding 37: Cooperation at Scale

The defining experiment ran 100, 250, 500, and 1,000 agents with spatial hashing, 4,000 ticks, five seeds each:

N	Survival	Clusters	Mean Cluster Size	Wall Time
100	64%	13.6	4.8	17s
250	49%	13.2	9.5	103s
500	74%	17.2	22.2	262s
1,000	96%	13.8	72.4	738s

Three observations:

Cooperation strengthens at scale. Survival increases from 64% at $N=100$ to 96% at $N=1,000$. More agents means more potential cooperation partners within harmony range, which means more resonance regeneration per agent. The cooperation benefit scales super-linearly with density.

Cluster count is stable. The number of clusters stays at approximately 14 regardless of population. This is a structural property of the arena: with 4+ energy wells distributed across a 400×400 space, the well geometry creates 3-4 natural basins of attraction, and agents within each basin form 3-4 sub-clusters within harmony range. The cluster count reflects arena geometry, not population.

Cluster size scales linearly. Mean cluster size is approximately $N/14$ — each cluster absorbs a proportional share of the population. There is no hard “Dunbar number” (a maximum group size beyond which clusters split). The limit is spatial, not cognitive: how many agents fit within harmony range of each other near a well.

The N=250 dip (49% survival) is the “scaling valley”: enough agents to create well competition, not enough density for universal cooperation. At N=100, agents are sparse enough that each finds a well without much competition. At N=1,000, agents are dense enough that cooperation partners are always nearby. At N=250, agents compete for wells but can’t always find partners — the worst of both regimes.

4.5 Determinism

All results in this book are deterministically reproducible. A lock-in test (`tests/determinism_lockin.rs`) runs a 500-tick simulation with 8 agents twice using the same seed and asserts bitwise identical final positions, energies, and cooperation counts. Both the lock-in test and a different-seeds-different-results test pass.

This means every experiment can be reproduced exactly:

```
\ExtensionTok{cargo} run \AttributeTok{{-}{-}example} cooperation\_1000 \AttributeTok{{-}{-}example}
```

The same seed produces the same output, every time, on the same platform. (Cross-platform determinism is not guaranteed due to floating-point implementation differences, but single-platform reproducibility is verified.)

Next: Chapter 5 — Cooperation as Thermodynamic Necessity

Part II

The Illusion

Chapter 5

A Note Before Part II

The following eight chapters present our original findings as we understood them before the controller replacement study of Chapter 13. The data is real. The observations are accurate. The interpretation was not.

We believed cooperation was thermodynamically inevitable — that the FEP gradient’s social component was the mechanism driving clustering and survival. We tested this belief against 45 environmental variations and it held every time. We were confident.

We were wrong.

We present these chapters unrevised because the experience of believing them is essential to understanding why the falsification in Part III matters. If we simply told you the correct answer, you would wonder why it took us so long. By showing you what we saw — the compelling patterns, the robust statistics, the consistent results — you will understand why the single experiment we did not think to run changed everything.

Read Part II as a detective story where the detective has the wrong suspect. The evidence is real. The clues are genuine. The conclusion is incorrect. And the moment of revelation, when it comes, is more powerful for having been delayed.

Chapter 6

Chapter 5: Cooperation as Thermodynamic Necessity

In which twenty digital agents, given nothing but energy budgets and the laws of physics, invent society.

6.1 The Simplest Possible World

Imagine a flat plane, two hundred units across. Scatter twenty agents at random. Give each one two hundred joules of energy and a simple rule: every tick, you lose a fraction of your energy to exist. When your energy reaches zero, you die.

Place two energy wells on the plane — fixed positions, finite capacity. An agent standing near a well regenerates energy slowly. But the wells deplete: each joule drawn is a joule gone.

Now add one more rule: if two agents are near each other and their harmony profiles resonate (a dot product above 0.5), both regenerate a small amount of energy. This is cooperation — not chosen, not communicated, not negotiated. It simply happens when compatible agents are proximate.

Finally, give each agent a behavioral gradient: a direction to move that reduces its estimated free energy. The gradient points toward energy wells when the agent is depleted, toward resonant partners when energy is moderate, and toward unexplored regions when energy is high. The agent follows this gradient. It has no memory, no learning, no communication, no strategy.

Press play.

6.2 What Happens

Within the first few hundred ticks, agents begin to cluster. Not at the energy wells — though some find wells early — but near each other. The FEP gradient's cooperation component draws agents toward resonant partners, and once a pair begins resonating, both gain energy, which reduces their depletion urgency, which allows them to stay together longer. A positive feedback loop forms: proximity → resonance → energy → reduced urgency → sustained proximity.

By tick 2,000, the twenty agents have organized into two or three groups, typically centered near the energy wells but not precisely on them. The groups are loose — agents drift in and out — but persistent. Agents that fail to join a group deplete their energy and collapse. Agents within groups survive.

By tick 5,000, with finite wells partially depleted, the dynamics become interesting. Groups that deplete their local well must migrate. The FEP gradient, sensing diminishing well returns, gradually shifts the group’s centroid toward the remaining well. This migration is not coordinated — no agent decides to move the group. Each agent independently follows its gradient, and the gradient happens to point the same direction for everyone in the group, because they all face the same depletion landscape.

By tick 10,000, a stable equilibrium has formed. The surviving agents (typically 60–80% of the original twenty) are clustered in one or two groups near the remaining energy sources, regenerating through a combination of well access and mutual resonance.

This is cooperation. It emerged from physics.

6.3 The Six-Condition Ablation

We did not accept this result at face value. Emergent cooperation in agent simulations has been reported many times, and the typical critique is that the cooperation was designed in — that the rules were tuned to produce the desired outcome.

We addressed this with a six-condition ablation study (Finding 2), removing components one at a time:

Condition	Clustering	Collapse Rate	Interpretation
FREE (no enforcement)	15.1	0%	No clustering — agents wander
ENERGY_ONLY	15.1	0%	Energy costs alone don’t cluster
E+OFFLOAD	15.1	0%	Epistemic offloading alone doesn’t cluster
E+GRADIENT	1.8	72%	Gradient clusters but kills
E+OFF+RANDOM	15.1	0%	Random gradient doesn’t cluster
FULL	13.5	45%	Sustainable clustering

The result is clean: the FEP gradient is the clustering mechanism (only conditions with the gradient produce clusters), and epistemic offloading (harmony resonance regeneration) is the survival mechanism (without it, 72% collapse). Neither alone is sufficient. Both together produce sustainable cooperation.

This decomposition was later confirmed by the null model experiment (Finding 32): random-walk agents (no FEP gradient) survive at 100% but don’t form social structure. Agents without resonance regeneration (NO_REGEN) form groups but die (3.9 out of 20 survive). The FEP gradient creates the structure; resonance provides the energy.

6.4 The Metric Independence Result

The most surprising finding from the ablation was not what we expected.

We had built the engine around Integrated Information Theory’s Φ metric — a measure of information integration inspired by Tononi (2004). We assumed Φ was important. So we tested five different consciousness metrics: IIT-inspired Φ , Shannon entropy, a random scalar, a constant value, and zero.

The first four produced statistically identical clustering: 4.57 ± 0.97 , 4.47 ± 1.02 , 4.28 ± 0.86 , and 4.57 ± 0.97 respectively (Finding 3). Twenty seeds each. No significant difference between any pair.

Only the $\Phi = 0$ condition differed: 7.50 ± 3.41 — clusters 35% looser than any nonzero metric (Finding 4).

This means the *specific* consciousness metric is irrelevant. IIT, Shannon entropy, even a constant — they all produce the same cooperation. What matters is having *some* nonzero integration variable that modulates the gradient's social component. The on/off distinction matters; the dial setting does not.

For the book's thesis, this is the universality result: cooperation depends on the thermodynamic cost structure, not the specific integration theory. It is analogous to Prigogine's dissipative structures (1977) — the macro-pattern emerges from the energy flow, independent of the microscopic details.

6.5 What We Could Not Kill

After establishing that cooperation emerges and identifying its mechanisms, we spent the next thirty experiments trying to destroy it.

We increased maintenance pressure from 0.05 to 1.50 joules per tick — a thirty-fold increase (Finding 38). Cooperation persisted.

We reduced harmony range from 40 to 10 units (Finding 38). Cooperation persisted.

We set ambient regeneration to zero (Finding 38). Cooperation persisted.

We fragmented perception so agents could only see 30 units (Finding 41). Cooperation fragmented into smaller groups but each group survived.

We added obstacles that pushed agents away from each other (Finding 41). Cooperation routed around them.

We moved the energy wells every 2,000 ticks (Finding 41). Cooperation found the new wells.

We removed all information asymmetry advantages (Finding 29). No effect.

We gave agents memory and learning (Findings 34, 31). No improvement.

We added communication — energy broadcasting and well location sharing (Finding 40). Zero effect.

We added conformal curvature that bent the geometry of the simulation space (Finding 36). Irrelevant.

Across forty-one findings and fifty-five experiments, we found exactly one way to kill cooperation: remove resonance regeneration entirely (Finding 32, NO_REGEN condition). Without mutual energy benefit, cooperation provides no survival advantage, and agents die alone.

6.6 The Implication

Cooperation in this engine is not a choice. It is not a strategy. It is not an emergent behavior of intelligence, memory, communication, or social learning. It is a thermodynamic inevitability — as certain as heat flowing from hot to cold, as predictable as water flowing downhill.

This does not mean cooperation is *easy*. The phase transition experiment (Chapter 6) shows that above a critical maintenance pressure, cooperation becomes bistable —

either everyone survives or everyone dies, depending on initial spatial configuration. The parameter sensitivity analysis (Finding 33) shows that harmony range and resonance regeneration rate are critical — reduce either enough and survival drops sharply.

But within the engine’s parameter space, cooperation is the default state. It requires no intelligence to produce and extraordinary intervention to prevent. The only thing that *can* prevent it is removing the mutual benefit mechanism itself — eliminating the reason agents benefit from proximity.

Or — as we discovered in Chapter 14 — giving agents the freedom to refuse.

Next: Chapter 6 — The Temperature of Society: Phase Transitions in Cooperative Systems

Chapter 7

Chapter 6: The Temperature of Society

In which we discover that cooperation does not decline gradually — it shatters.

7.1 The Ising Analogy

In 1925, Ernst Ising solved a one-dimensional model of ferromagnetism. Below a critical temperature, atomic spins align spontaneously — the material becomes magnetic. Above it, thermal noise overwhelms the alignment force and the material becomes paramagnetic. The transition is sharp. There is no gradual demagnetization. One degree below the critical temperature, the magnet works. One degree above, it does not.

We found the same pattern in cooperation.

7.2 The Pressure Sweep

We swept maintenance pressure — the energy cost of existing — from 0.05 to 1.00 joules per tick across twelve levels, with eight seeds at each level (Finding 20). Maintenance pressure is our analogue of temperature: higher pressure means more thermal noise in the system, more entropy production, more urgency.

The result was not a gradual decline.

Pressure	Survival	Cooperation Events	Temperature	Entropy
0.05	100%	378K	311.5K	66.6 J/K
0.10	60%	507K	316.6K	201.7 J/K
0.20	40%	1.07M	330.1K	523.5 J/K
0.30	27%	1.19M	331.4K	558.4 J/K
0.75	13%	1.30M	330.7K	571.2 J/K
1.00	38%	1.31M	330.6K	571.6 J/K

At 0.05 J/tick, every agent survives. At 0.20 J/tick, only 40% survive. The transition happens in a narrow band — a four-fold increase in pressure causes a 60% drop in survival. This is not linear degradation. It is a phase transition.

7.3 The Bistability

The most telling feature of the data above 0.20 J/tick was not the average survival but the variance. At pressures above the critical threshold, individual seeds did not produce moderate survival (10 out of 20). They produced either perfect survival (20 out of 20) or total collapse (0 out of 20). The distribution was bimodal.

This is the hallmark of a first-order phase transition. The system has two stable states — cooperative (all survive) and collapsed (all die) — and the initial condition determines which state the system falls into. Small differences in initial spatial configuration — who happened to start near a well, who happened to start near a compatible partner — determined whether the entire population lived or died.

In the Ising model, this corresponds to the coexistence region where both phases are thermodynamically stable. A small perturbation — a nucleation event — tips the system one way or the other. In our engine, the “nucleation event” is whether the first few agents form a viable cooperative cluster before their energy depletes. If they do, the cluster grows and absorbs nearby agents. If they don’t, everyone dies alone.

7.4 Desperation Cooperation

The counterintuitive finding was in the cooperation column. At low pressure (0.05 J/tick), agents cooperated 378,000 times across the simulation. At high pressure (1.00 J/tick), they cooperated 1.31 million times — a 3.5-fold increase.

Agents cooperate *more* when they are dying.

This is not paradoxical once you understand the FEP gradient. The cooperation component of the gradient is weighted by $(1.0 - \text{energy_fraction})$ — urgency increases as energy drops. Low-energy agents seek partners more desperately than comfortable ones. Under high pressure, every agent is low-energy, so every agent is urgently seeking cooperation. The result is frantic mutual aid — cooperation born not from abundance but from desperation.

This maps to a documented sociological phenomenon. Fritz (1996) studied community response to disasters and found that social cohesion *increases* immediately after catastrophe. Neighbors who never spoke before share food, shelter, and labor. The “therapeutic community” that forms in the aftermath of disaster is a real-world example of desperation cooperation — and it eventually dissipates as conditions normalize, just as our high-pressure cooperation events decrease per capita as agents collapse.

7.5 The Thermodynamic Response

The engine tracks physical thermodynamic quantities for each agent: temperature, entropy, and Helmholtz free energy ($F = U - TS$).

Across the pressure sweep, system temperature rose monotonically from 311.5K to 330.6K — a 19K increase driven entirely by energy dissipation from maintenance costs and collision heat. Entropy production increased 8.6-fold from 66.6 to 571.6 J/K. And Helmholtz free energy — the maximum extractable work — dropped to zero above the critical pressure.

Zero available work means the system has no thermodynamic capacity for ordered behavior. Every joule of internal energy is locked up as entropy. The agents are alive (some

of them) but thermodynamically exhausted — their energy goes entirely to maintenance with nothing left for directed action.

This is the thermodynamic definition of the collapsed phase: a system at maximum entropy production with zero free energy. It exists, but it cannot organize.

7.6 The Parameter Sensitivity

The phase transition raised an obvious question: is the critical pressure at 0.20 J/tick a universal feature of the engine or an artifact of the specific parameter values in `ThermodynamicConstants`?

The parameter sensitivity analysis (Finding 33) addressed this directly. We swept each of four parameters independently by $\pm 50\%$:

Parameter	Max Effect Size	Verdict
Maintenance per tick	$d = 1.13$	CRITICAL
Harmony range	$d = 1.29$	CRITICAL
Resonance regen rate	$d = 0.63$	Moderate
Ambient regen rate	$d = 0.26$	Robust

Maintenance pressure and harmony range are the critical parameters — unsurprisingly, since they directly control the two mechanisms identified in Chapter 5 (gradient urgency and cooperation range). Ambient regeneration has almost no effect ($d = 0.26$), confirming that slow background recovery is irrelevant compared to cooperation.

The 2^3 factorial design (Finding 35) went further, testing all eight combinations of high/low for the three most important parameters. The interaction between harmony range and resonance regeneration rate was synergistic (+3.32 agents), meaning their combined effect was greater than the sum of their individual effects. Wide social range $\times$ high mutual benefit = a multiplicative cooperation advantage.

The best parameter combination (low maintenance, high range, high regen) produced 19.9 out of 20 survival. The worst (high maintenance, low range, low regen) produced 9.2 out of 20. That is a 2.2-fold difference in survival from parameter choice alone — confirming that the phase transition is real and parameter-dependent, not an artifact of a single tuning.

7.7 The Critical Infrastructure

The factorial result has a direct policy interpretation, though we offer it cautiously given the engine’s limitations.

If you want to maximize cooperation in a thermodynamic agent system, you should invest in social range (the distance over which agents can interact with each other) and mutual benefit rate (how much agents gain from being near compatible partners). Reducing maintenance pressure (the cost of existing) helps, but less than expanding social infrastructure.

In human terms: reducing the cost of living matters, but building community centers matters more. The synergistic interaction means that neither alone is sufficient — you need both proximity and benefit. A community center in a neighborhood where people can’t afford to leave their homes (high maintenance) produces less cooperation than the same center in a neighborhood with moderate costs.

This is Klinenberg’s thesis in *Palaces for the People* (2018): social infrastructure — the physical spaces and institutions that enable social contact — is the critical variable for community resilience. Our engine, knowing nothing about Klinenberg, arrives at the same conclusion from thermodynamics.

7.8 The Phase Diagram

Combining the pressure sweep, the parameter sensitivity, and the factorial design, we can sketch a phase diagram for cooperation:

Cooperative phase (below critical pressure): All agents survive. Cooperation is thermodynamically favorable — the energy gained from resonance exceeds the energy spent on maintenance. The system reaches a stable equilibrium with organized social structure.

Transition zone (near critical pressure): Bistable. The outcome depends on initial conditions. Early cooperative clusters nucleate and grow; isolated agents collapse. Small perturbations determine the fate of the entire population.

Collapsed phase (above critical pressure): Most agents die. Those that survive do so near wells, not through cooperation. Entropy production is maximal, free energy is zero. The system exists but cannot organize.

The critical pressure is not a fixed value — it shifts with harmony range (wider range lowers the critical pressure, making cooperation viable under harsher conditions) and resonance regeneration rate (higher regen provides more energy per cooperative interaction, buffering against pressure).

This is the thermodynamics of togetherness in a single diagram: a phase space where the axes are *pressure* (cost of existing) and *infrastructure* (range $\times$ benefit), and the boundary between cooperation and collapse is a curve, not a point.

Next: Chapter 7 — Social Bonds and the Physics of Friendship

Chapter 8

Chapter 7: Social Bonds and the Physics of Friendship

In which physical constraints between resonant agents improve survival by 70%, and we discover that friendship is a spring.

8.1 The Constraint Hypothesis

Chapters 5 and 6 established that agents cluster and cooperate through the FEP gradient. But the gradient only *attracts* — it creates no lasting connection. Once an agent’s energy rises above 50%, the well-seeking urgency drops, and the social component weakens. Agents drift apart.

Real social bonds don’t work this way. Friendships persist through periods of comfort. Family ties endure when no one needs help. The question is: does adding *persistence* to social connections — physical constraints that hold resonant agents together — change cooperation outcomes?

We modeled this with the simplest possible mechanism: a distance constraint (a spring) between agents whose harmony resonance exceeds 0.7. Once formed, bonds persist — the engine has no remove-constraint API, modeling the sociological observation that strong social bonds are hard to break once formed.

8.2 Finding 19: Bonds Transform Survival

Twenty agents, six distinct harmony archetypes (ensuring selective bonding, not universal), ten seeds each. BONDS condition forms springs between high-resonance pairs. NO_BONDS condition is the standard FEP-only baseline.

Metric	NO BONDS	BONDS	Effect
Survival	11.4/20 (57%)	19.4/20 (97%)	p=0.001, d=2.44
Energy	75.5 J	99.1 J	p=0.019
Bonds formed	0	33.6/190 (17.7%)	Selective

A 70% survival improvement. Cohen’s d of 2.44 — the largest effect size of any experiment in this book.

The bonds were selective: only 33.6 out of 190 possible pairs (17.7%) formed connections. The six harmony archetypes created a compatibility landscape where some pairs resonated above 0.7 and others didn’t. The bonds that formed were genuine resonance matches, not universal connectivity.

The mechanism: bonds hold agents near resonant partners even when the FEP gradient’s social component weakens (which happens as energy rises above 50%). Without bonds, comfortable agents drift apart and may not reconnect before the next energy crisis. With bonds, they stay together — maintaining a constant trickle of resonance regeneration that compounds over thousands of ticks into a massive survival advantage.

Friendship, in this engine, is a spring that holds you near the people who keep you alive — even when you forget, in moments of comfort, that you need them.

8.3 The Dunbar Limit

The large-scale experiment (Finding 37, Chapter 4) showed cooperation scaling to $N=1,000$. But an earlier experiment (Finding 21-23) tested intermediate populations to look for natural group size limits — the engine’s analogue of Dunbar’s number.

Population	Survival	Clusters	Mean Cluster Size
4	70%	2.5	1.1
12	90%	3.3	3.0
20	70%	4.0	4.0
48	50%	3.2	12.8
96	60%	4.2	14.7

Optimal population: $N=12$. Survival peaks at 90%, then declines as more agents compete for the same fixed wells. The number of clusters stays remarkably stable at 3-4 regardless of population — the arena’s well placement creates natural basins of attraction.

Mean cluster size grows linearly with population (approximately $N/4$), suggesting no hard Dunbar limit. The constraint is not cognitive (our agents have no cognition) but spatial: harmony range (40 units) defines the maximum interaction distance, and the arena geometry determines how many agents can fit within that range of a well.

At $N=1,000$ (Chapter 4), cluster size reached 72.4 with 14 clusters — still stable. The “Dunbar number” of this engine is not a number but a *density*: how many agents per unit area can sustain resonance within harmony range. This scales with the arena, not with the population.

8.4 Migration

The migration experiment (Finding 27) placed all agents near the first of four staggered wells (capacities: 800, 1500, 2500, 4000 joules). As the first well depleted, agents needed to discover and migrate to the others.

Result: 18.2 out of 20 agents migrated, traveling a mean distance of 3,274 units — more than thirty times the arena width. All twenty survived. The FEP gradient,

sensing declining well returns, redirected agents toward the remaining wells without any coordination or communication.

The migration was not a collective decision. Each agent independently followed its gradient, which happened to point in the same direction because every agent faced the same depletion landscape. Collective migration emerged from individual physics.

8.5 Generational Wealth

Finding 28 tested whether position advantage persists across generations. In an evolutionary simulation (20 agents, 20,000 ticks, reproduction every 1,000 ticks), we tracked whether agents born near wells — the “trust fund babies” of our simulation — dominated the population over 19 generations.

They did not. Twenty-five percent of the final population descended from the initial well-adjacent quartile — exactly the random expectation. Lineage diversity remained at 20 (all original lines survived). The Gini coefficient stayed stable at 0.24-0.25 across generations.

Cooperation erased hereditary advantage. The resonance regeneration mechanism distributed energy to all nearby agents regardless of lineage, and the FEP gradient brought agents together regardless of origin. Starting near a well gave a brief advantage, but evolution + cooperation equalized outcomes within a few generations.

This finding contrasts sharply with the equal opportunity paradox of Chapter 9, where identical starting conditions produced maximum inequality. The difference: the inequality experiment had no evolution. In a single generation, small advantages compound. Across multiple generations with reproduction, the advantages wash out through cooperative energy sharing and random mutation.

The policy implication — again offered cautiously — is that *mobility* (the ability to form new cooperative relationships) matters more than *inheritance* (starting position). An engine with both mobility and cooperation produces equality. An engine with neither produces the Matthew Effect.

Next: Chapter 8 — Raiders, Evolvers, and the Solidarity Response

Chapter 9

Chapter 8: Raiders, Evolvers, and the Solidarity Response

In which adversaries die faster than cooperators, societies evolve universal compatibility, and the arms race favors the social.

9.1 The Adversarial Threshold

Every utopian system faces the same question: what happens when someone defects?

We introduced adversarial agents — entities with inverted harmony profiles that produce low resonance (below 0.2) with cooperators. At low resonance, interactions *drain* energy rather than regenerating it. Adversaries are energetic parasites: their presence near cooperators costs energy without reciprocal benefit.

We swept the adversary fraction from 0% to 50% across nine levels, ten seeds each (Finding 14).

The result was a threshold function. Below 20% adversaries, cooperator survival remained above 80%. At 25%, it dropped sharply. Above 30%, cooperator societies collapsed.

But the adversaries fared worse. At every fraction tested, adversaries died faster than cooperators (Finding 15). Their inverted harmony profiles meant they couldn't regenerate from *any* interaction — not with cooperators (dissonant) and not with each other (all adversaries share the same inverted profile, so same-type resonance is low). The adversarial strategy is energetically self-defeating.

This is the thermodynamic argument against free-riding: defection costs energy. In a game theory framework, defection is free — the defector simply takes the cooperator's contribution without paying. In a thermodynamic framework, defection means proximity without resonance, which means collision energy drain without regeneration. The defector pays a physical price for being antisocial.

9.2 The Solidarity Response

What do cooperators do when adversaries are present?

The evolutionary experiment (Finding 18) ran 16 cooperators against 4 fixed adversaries for 30,000 ticks with reproduction every 1,000 ticks. Dead cooperators were replaced by offspring of the fittest survivor, with small mutations to harmony profiles.

We expected tribalism. We expected cooperators to evolve narrow harmony profiles — specializing to resonate only with their own type, excluding outsiders and adversaries alike.

Instead, cooperators evolved *universal compatibility*. Mean pairwise resonance increased from 0.86 to 0.996 — approaching the theoretical maximum of 1.0. Cooperators became compatible with *everyone*, not just their kin.

The logic is thermodynamic. An agent with a narrow harmony profile can resonate with few partners. An agent with a broad profile can resonate with many. Under scarcity, the agent that can cooperate with anyone is more likely to find a partner and survive. Natural selection, operating on energy fitness, selects for social breadth rather than social depth.

This is the opposite of the tribalism prediction from evolutionary psychology. Our engine suggests that adversarial pressure drives *openness*, not *closure*. Societies under threat become more inclusive, not less — because exclusivity is an energy luxury that dying agents cannot afford.

9.3 The Arms Race

The adversarial learning experiment (Finding 39) asked: what happens when both sides can adapt?

Three conditions: fixed cooperators vs fixed adversaries (baseline), learning cooperators vs fixed adversaries, and learning cooperators vs learning adversaries. “Learning” here means reward-modulated plasticity on the FEP gradient weights — agents adjust how much they weight cooperation, well-seeking, and danger avoidance based on their energy trajectory over the last 100 ticks.

Condition	Coop Alive	Adv Alive	Coop Weight Drift	Adv Weight Drift
Fixed vs Fixed	6.7/15	2.2/5	0.00	0.00
Learn vs Fixed	5.5/15	1.6/5	2.82	0.00
Learn vs Learn	9.8/15	2.9/5	2.62	3.06

When only cooperators learn, they slightly *underperform* the fixed baseline (5.5 vs 6.7). Learning introduces instability into a system that was already at a local optimum (Finding 31 showed the default weights are near-optimal). One-sided adaptation is worse than no adaptation.

But when both sides learn, cooperators surge to 9.8/15 — a 46% improvement over the fixed baseline. The arms race *favors cooperation*. Adversary weight drift (3.06) is higher than cooperator drift (2.62), suggesting adversaries need to adapt more aggressively to survive, but their adaptations are less effective because their fundamental strategy (parasitism) is energetically unviable.

The result was not statistically significant after Holm-Bonferroni correction ($d = -0.46$, approaching medium), but the trend is consistent across seeds and the mechanism is clear: mutual adaptation selects for the strategy with more degrees of freedom, and cooperation has more degrees of freedom than parasitism. A cooperator can adjust which

partners to seek, how urgently to cluster, and how much to weight well-seeking vs social seeking. A parasite can only adjust how aggressively to pursue victims — but pursuing victims harder doesn't fix the resonance problem.

9.4 The Evolutionary Lesson

Across findings 14-18 and 39, a consistent picture emerges:

1. **Adversaries are self-defeating:** their strategy costs energy in a thermodynamic system
2. **Cooperators evolve solidarity, not tribalism:** adversarial pressure broadens social compatibility
3. **Unilateral adaptation hurts:** changing strategy in a stable environment introduces instability
4. **Mutual adaptation favors cooperators:** the arms race selects for the more flexible strategy

This is not the prisoner's dilemma. In the prisoner's dilemma, defection is free and cooperation is costly. In thermodynamic cooperation, defection is costly (proximity without resonance drains energy) and cooperation is free (proximity with resonance generates energy). The payoff matrix is inverted: the "temptation to defect" is negative, because defection actively harms the defector.

No wonder cooperation is robust. The engine doesn't just reward cooperation — it punishes defection with physics.

Next: Chapter 9 — The Economics of Scarcity

Chapter 10

Chapter 9: The Economics of Scarcity

In which abundance kills, charity fails, and identical starting conditions produce maximum inequality.

10.1 Three Parables

This chapter presents three economic experiments that each contradict a common intuition. Together, they form a thermodynamic theory of resource distribution that is more pessimistic than Marx and more optimistic than Malthus.

10.1.1 Parable One: The Curse of Abundance

We placed twenty-four agents in an arena with three energy wells under three conditions: abundant wells (10,000 joules each), scarce wells (1,000 joules each), and depleting wells (2,000 joules with a crowding penalty — wells drain faster when more agents stand on them).

Common sense predicts that abundant resources produce the best outcomes. Common sense is wrong.

Condition	Survival	First Well Depleted
ABUNDANT (10K each)	7.8/24	tick 8,699
SCARCE (1K each)	10.0/24	tick 1,611
DEPLETING (2K, crowd penalty)	13.9/24	tick 1,387

Abundant agents survived worst. Depleting agents — facing the harshest resource dynamics — survived best.

The mechanism is cooperation timing. Abundant agents have no urgency. Their energy stays high, so the FEP gradient's cooperation component (which scales with `1.0 - energy_fraction`) remains weak. They wander. They don't cluster. They don't cooperate. Then, around tick 8,000, the abundant wells finally deplete, and agents who never learned to cooperate find themselves alone, energy-depleted, with no partners nearby.

Scarce agents face immediate pressure. Their wells deplete by tick 1,600, forcing early migration and cooperation. They cluster quickly, establish resonance partnerships, and sustain each other through mutual energy regeneration. The scarcity is the catalyst.

Depleting agents face the most interesting dynamics. The crowding penalty means standing at a well with too many neighbors drains the well faster — but each agent only receives the base regeneration rate. The well’s energy is wasted on crowd overhead. This forces agents to *disperse* across multiple wells rather than crowding a single one — an emergent resource distribution strategy that no agent planned.

Economists recognize this pattern. The “resource curse” (Sachs & Warner 1995) documents how nations with abundant natural resources often experience worse economic outcomes than resource-poor nations — a phenomenon attributed to institutional decay under conditions of easy wealth. Our engine reproduces the resource curse from thermodynamics: abundance removes the pressure that drives cooperation, and cooperation is the only sustainable energy source.

10.1.2 Parable Two: The Futility of Charity

We tested whether agents could rescue their collapsed neighbors (Finding 25). Three conditions: no rescue, free rescue (collapsed agents automatically revived when a living agent is nearby), and costly rescue (the rescuer donates 20% of their energy to revive the collapsed agent with 30 joules).

Condition	Survival	Rescues	Re-collapses
NO RESCUE	10.2/20	0	0
FREE RESCUE	10.4/20	40	40 (100%)
COSTLY RESCUE	10.2/20	3	3 (100%)

Every rescued agent re-collapsed. One hundred percent recidivism.

The gift of 30 joules is consumed by maintenance faster than the rescued agent can reach a well or find a resonant partner. The agent revives, burns through its gift in approximately 120 ticks, and collapses again. The donor, meanwhile, has lost 20% of its energy — weakening itself without producing any lasting benefit.

This is not a failure of generosity. It is a structural impossibility. Individual energy transfer cannot substitute for *access to sustainable energy sources* (wells and resonant partners). A collapsed agent needs to be near a well AND near a compatible partner to sustain itself. Receiving a one-time energy gift addresses neither requirement.

The parallel to development economics is direct. Easterly (2006) and Moyo (2009) documented how foreign aid without institutional support — schools without teachers, hospitals without supply chains, wells without maintenance — fails to produce lasting improvement. Our engine reproduces this finding: energy transfer without structural change is futile. The rescued agent needs infrastructure (proximity to wells and partners), not charity (one-time energy).

10.1.3 Parable Three: The Equal Opportunity Paradox

This is the most unsettling result.

We started all twenty-four agents at the center of the arena with identical energy (200 joules each) under three conditions: equal start (everyone at center), random start (random positions, same energy), and unequal start (random positions, energy proportional to well proximity).

Condition	Starting Gini	Final Gini	Interpretation
EQUAL	0.000	0.657	Maximum inequality from equality
RANDOM	0.000	0.348	Moderate inequality
UNEQUAL	0.081	0.365	Less inequality than equal start

Equal starting conditions produced the highest Gini coefficient — 0.657, comparable to South Africa (0.63), the most unequal major economy on Earth.

The mechanism is symmetry breaking. When all agents start at the same position with the same energy, the FEP gradient points them all in the same direction — toward the nearest well. They race. The winners (agents whose stochastic position perturbation placed them slightly closer to the well) arrive first and begin regenerating. The losers arrive later, find the well crowded, and begin depleting. Small initial advantages — fractions of a unit of distance — compound through positive feedback: energy → reduced urgency → sustained well access → more energy.

This is the Matthew Effect (Merton 1968): “For unto every one that hath shall be given, and he shall have abundance: but from him that hath not shall be taken away even that which he hath.” Our engine derives the Matthew Effect from thermodynamics, without any social mechanism, institutional bias, or wealth inheritance. Pure physics produces maximum inequality from perfect equality.

Random starting positions produce *less* inequality (Gini 0.348) because the randomness distributes agents across the arena. Some start near wells, others near partners — but the spatial diversity means no single well faces a competitive crush. The inequality of starting position paradoxically produces equality of outcome.

Mobility was highest in the equal condition (95.8% of agents changed quartile) — not because the society was fair, but because everyone’s fate was determined by the race to the well, not by their starting position. High mobility + high inequality is the signature of a winner-take-all system.

10.2 The Unified Theory

These three parables converge on a single principle:

Cooperation is the only sustainable energy source, and anything that substitutes for cooperation — abundant resources, individual charity, equal opportunity — ultimately undermines it.

Abundance removes the pressure to cooperate. Charity transfers energy without creating cooperative structure. Equal starting conditions create competition rather than complementarity.

What works is scarcity (which creates cooperation pressure), infrastructure (which enables cooperation), and diversity (which ensures agents find complementary partners rather than identical competitors).

This is not an argument for imposed scarcity or against aid. It is a thermodynamic observation: systems that rely on cooperation for survival are resilient. Systems that

rely on individual resource access are fragile. The design of the energy landscape determines which regime the system enters — and well-meaning interventions that make individual survival easier can inadvertently undermine the cooperative structures that make collective survival possible.

Next: Chapter 10 — Technology and the Shrinking of Range

Chapter 11

Chapter 10: Technology and the Shrinking of Range

In which the internet helps and algorithms hurt, and we learn that connection is not the same as proximity.

11.1 Five Eras

We modeled the evolution of communication technology as changes to the FEP gradient's parameters across five eras (Finding 12-13), each representing a different topology of social interaction:

Pre-Internet (1970s): Small harmony range (25 units), high resonance regen, no wells beyond local. Agents can only interact with those physically nearby, but those interactions are rich. This models the neighborhood: you know your neighbors, you don't know the world.

Early Internet (1990s): Medium range (40 units), moderate regen. Agents can sense more partners but each connection is weaker. Email and early web: broader reach, shallower engagement.

Social Media (2010s): Large range (60 units), lower regen per interaction. Agents see many potential partners but gain little from each. Facebook and Twitter: vast networks, thin connections.

Algorithm Era (2020s): Very large range (80 units), very low regen, with “filter bubble” effect — agents only perceive partners with resonance above 0.6 (low-resonance agents are hidden by the algorithm). The algorithm maximizes the appearance of connection while reducing actual cooperative diversity.

Post-Algorithm (future): Maximum range (100 units), minimal regen, no filtering. Pure information: agents see everything and everyone but benefit from almost nobody.

11.2 The Results

Pre-internet agents — with the smallest range but the highest per-interaction benefit — produced the most cooperation. Algorithm-era agents produced the least.

The mechanism is the range-regen tradeoff identified in Chapter 6. Social range and resonance regeneration rate interact synergistically (Finding 35, +3.32 agents). Increasing

range while decreasing regen moves along the *wrong* diagonal of the interaction surface — you lose more from weakened connections than you gain from broader reach.

The algorithm-era filter bubble compounds this: by hiding low-resonance agents, the algorithm reduces the probability of serendipitous encounters with partners who might have moderate (but positive) resonance. The agents who remain visible are the high-resonance matches — but the algorithm can’t increase the resonance *benefit*, only the visibility. The result is a smaller effective population with the same weak connections.

11.3 Connection Is Not Proximity

The deepest insight from this experiment is not about any specific technology era. It is about the distinction between *connection* (the ability to perceive others) and *proximity* (the ability to benefit from others).

In the engine, perception range (how far an agent can see) is different from harmony range (how close agents must be to resonate). Increasing perception range allows agents to detect distant partners and navigate toward them. But the benefit only occurs within harmony range — the physical distance at which resonance regenerates energy.

Technology, in our model, increases perception range but does not increase harmony range. You can see more people, but you still need to be *near* them to cooperate. And “near” is defined by physics, not by fiber optics.

The pre-internet era had the smallest perception range but agents were already within harmony range of their neighbors (because the arena was small relative to social range). Everyone they could see, they could cooperate with. Connection and proximity were identical.

The algorithm era has the largest perception range, but agents are distributed across a vast arena. They can see partners at 80 units but can only cooperate at 40. They spend energy traveling to perceived partners, arrive, cooperate briefly, then see a “better” match at 80 units and travel again. The restlessness of infinite choice degrades the sustained proximity that cooperation requires.

This maps directly to Turkle’s *Alone Together* (2011): technology that increases connection while decreasing proximity produces loneliness. Our engine produces the same pattern from thermodynamics: agents with maximal perception and minimal benefit are the loneliest — and the most likely to die.

Next: Chapter 11 — Learning and the Paradox of Optimal Defaults

Chapter 12

Chapter 11: Learning and the Paradox of Optimal Defaults

In which agents that learn discover they already knew everything, and the couch potato strategy turns out to be immortal.

12.1 The Learning Hypothesis

If cooperation is thermodynamically inevitable (Chapter 5), a natural question follows: can agents learn to cooperate *better*? Can reward-modulated plasticity — agents adjusting their behavioral weights based on energy outcomes — improve on the hardcoded defaults?

We built the infrastructure to test this. Each agent received four learnable weights controlling how strongly the FEP gradient weighted cooperation, well-seeking, danger avoidance, and exploration. A reward signal (the average energy trend over the last 100 ticks) drove weight updates via momentum-stabilized gradient ascent. This is not Hebbian learning (which uses local pre/post-synaptic correlations) but reward-modulated plasticity — a form of reinforcement learning where a global scalar reward shapes behavioral parameters.

The answer to the learning hypothesis is: no. And the reason is more interesting than the answer.

12.2 Finding 31: Weights Are Already Optimal

Twenty seeds, 10,000 ticks, two conditions: fixed weights (baseline) and learning weights (reward-modulated plasticity). Result:

Metric	Fixed	Learning	Effect
Survival	11.4/20	7.5/20	d=0.43 (n.s.)
Cooperation	1.59M	1.54M	d=0.15 (n.s.)
Weight drift	0.000	0.018	—
Learned cooperation weight	2.000	1.989	-0.5% drift
Learned well weight	4.000	3.994	-0.2% drift

Metric	Fixed	Learning	Effect
Learned danger weight	3.000	3.000	0% drift

The weights barely moved. In 10,000 ticks of reward-modulated plasticity with momentum, the cooperation weight drifted from 2.000 to 1.989 — a 0.5% change. The danger weight didn’t move at all.

The per-tick reward signal (energy fraction delta) was too noisy to drive meaningful learning. Each tick’s energy change is dominated by the constant maintenance cost, with cooperation and well regeneration adding small perturbations. The signal-to-noise ratio was insufficient for gradient-based learning.

12.3 Finding 34: Memory Doesn’t Help Either

We gave agents persistent memory — discovered well locations and partner resonance history — plus a windowed reward signal (100-tick energy trend average, 5× stronger learning rate). Agents could navigate to remembered wells and preferentially seek high-resonance partners.

Three conditions, twenty seeds each:

Condition	Survival	Weight Drift
Stateless	15.9/20	0.00
Memory only	13.9/20	0.00
Memory + learning	13.2/20	1.38

The windowed reward produced meaningful weight drift (1.38 vs 0.018) — the stronger signal worked as intended. But the result was *worse*, not better. Stateless agents survived at 15.9. Memory+learning agents survived at 13.2.

Memory agents discovered fewer than 1 well on average. Partner history filled up (19 of 20 max entries) but didn’t improve partner selection. The FEP gradient already pointed agents toward compatible partners — remembering which partners were good didn’t change which direction the gradient pointed.

The paradox: agents with more information and more learning capacity performed worse than agents with neither. The additional complexity introduced behavioral instability (weight drift disrupts a system already at local optimum) without providing usable information (the gradient already encodes the relevant environmental state).

12.4 Finding 44: The U-Curve Mechanism

The volitional cooperation experiment (Chapter 14) revealed a U-shaped survival curve. We diagnosed its mechanism through per-tick energy flow tracking:

Willingness	Well Regen	Resonance Regen	Efficiency	At Well%
1.0 (full coop)	4,979J	83,942J	2.75	44.7%
0.6 (partial)	5,000J	69,717J	2.31	44.1%
0.0 (selfish)	1,567J	22,007J	0.72	8.2%

Partial cooperators ($w=0.6$) get the same well energy as full cooperators (5,000J vs 4,979J) but 17% less resonance (70K vs 84K). They're distracted by the social gradient 60% of the time — sometimes moving toward partners, sometimes ignoring them — but their inconsistency means they cluster less tightly and receive less resonance per tick.

Full defectors ($w=0.0$) have the worst energy efficiency (0.72 vs 2.75) but the highest survival (20/20). How? They barely move. Displacement is 0.14 per tick (vs 0.17 for cooperators). They spend only 8.2% of ticks at wells (vs 44.7%). They sit still, spend almost nothing on movement, and survive by consuming minimal energy. When another agent happens to wander nearby, they passively benefit from resonance (22K events) without having sought it.

This is the couch potato strategy: do nothing, spend nothing, survive by inaction. It works because the engine's maintenance cost is per-tick, not per-action. An agent that stands still pays the same maintenance as one that traverses the arena — but the still agent conserves the kinetic energy that movement would have cost.

12.5 Finding 45: Harmony Dimensionality Is Irrelevant

We used eight harmony channels throughout the book. Was this necessary?

We swept active channels from 1 to 8 by zeroing out higher channels:

Active Channels	Survival	Mean Resonance	Resonance Variance
1	11.4	1.000	0.000
2	11.9	0.893	0.013
4	8.3	0.798	0.016
8	11.8	0.824	0.013

With one channel, everyone resonates with everyone (mean = 1.0, variance = 0.0). There is no social differentiation — every agent is equally compatible with every other. Survival (11.4) is essentially identical to eight channels (11.8), Cohen's $d = -0.055$.

More channels add social nuance (variance increases from 0.0 to 0.013) but don't change outcomes. The eight harmonies create a richer social landscape — agents have graded compatibility rather than binary match/mismatch — but this grading doesn't improve cooperation. It just makes it more selective.

The engine's cooperation doesn't need eight dimensions of social compatibility. It needs one: are you nearby and alive?

12.6 The Paradox of Optimal Defaults

Across four findings, a single pattern: the engine's default configuration is already near-optimal for cooperation, and additional complexity (learning, memory, communication, more harmony dimensions) either has no effect or makes things worse.

This is not a failure of the additions. It is a consequence of the FEP gradient's effectiveness. The gradient solves the cooperation problem so completely that there is nothing left for learning, memory, or communication to improve. Every relevant piece of information (where are the wells? where are compatible partners? how depleted am

I?) is already encoded in the gradient's four components. Adding redundant information sources adds noise without signal.

The implication for the book's thesis: cooperation in this engine is not just thermodynamically inevitable — it is *informationally simple*. It requires exactly one behavioral rule (follow the free energy gradient) and exactly one physical mechanism (harmony resonance regeneration). Everything else we tried to add was, in the strict thermodynamic sense, waste heat.

Next: Chapter 12 — Higher Dimensions and the Geometry of Cooperation

Chapter 13

Chapter 12: Higher Dimensions and the Geometry of Cooperation

In which cooperation survives the third dimension, curvature proves beautiful but useless, and we learn that social physics is dimension-independent.

13.1 The 3D Test

Every experiment in this book runs in 2D. The engine supports 2D through 9D — the physics, the harmony fields, the FEP gradient, and the spatial hash are all dimension-generic (using Rust’s `const D: usize` generics). But claiming N-dimensional support without testing it is an empty claim.

Finding 30 compared identical experiments in 2D and 3D: 20 agents, 8,000 ticks, 8 seeds, same harmony profiles, same thermodynamic constants.

Dimension	Survival	Clustering	Cooperation
2D	15.2/20	5.77	956K
3D	13.4/20	14.71	617K

Survival is not significantly different ($p=0.32$, $d=0.27$). Cooperation drops 35% ($d=1.11$, large effect). Clustering loosens $2.5\times$.

The mechanism is volume. A 100×100 2D arena has 10,000 square units. A $100\times 100\times 100$ 3D arena has 1,000,000 cubic units — $100\times$ the volume. Agents are $100\times$ more dilute. The harmony range (40 units) covers a sphere of volume 268,000 cubic units in 3D vs a circle of 5,027 square units in 2D — the 3D interaction volume is $53\times$ larger, but the arena volume grew $100\times$, so the probability of a random agent being within harmony range is roughly halved.

Agents cooperate less in 3D because they encounter each other less frequently. But the FEP gradient still drives them toward wells and partners, and the cooperation that does occur is still sufficient for survival. The social physics is dimension-independent; the social geometry is not.

13.2 Curvature: Beautiful Mathematics, Zero Social Effect

The conformal curvature system (Chapter 2, Channel 5) implements correct Riemannian geometry. The conformal factor $\sigma(x) = \kappa \times E_harmony(x)$ curves the simulation space near high-harmony regions. The geodesic correction:

$$a = -2(v \cdot \nabla \sigma)v + |v|^2 \nabla \sigma$$

produces measurable trajectory deflection. Finding 8 demonstrated this: a test body launched past a stationary high- Φ source deflects proportionally to the curvature scale κ , with Mann-Whitney $U = 0.0$, $p = 0.009$, Cohen's $d = -35.5$. The deflection is real, large, and statistically overwhelming.

Finding 36 tested whether this deflection improves cooperation. Three conditions: flat ($\kappa=0$), low curvature ($\kappa=0.01$), and high curvature ($\kappa=0.05$). Result: $d = 0.006$ for high curvature vs flat. Survival: 11.7 vs 11.7. Clustering: 6.41 vs 6.90. Cooperation events: identical within noise.

The curvature creates “harmony wells” in the metric — regions where space is compressed, making geodesics curve toward high-harmony areas. But the FEP gradient already drives agents toward these same areas. The curvature adds a second-order correction to trajectories that are already pointing in the right direction. It's like adding a gentle breeze to a river current — technically real, practically invisible.

This is an honest negative result. We invested 85 lines of tensor calculus in a system that produces the strongest p-value in the book (0.009) for trajectory deflection but contributes nothing to the social dynamics that the book is about. The mathematics is impeccable. The social relevance is zero.

For the GROUNDED/INSPIRED/SPECULATIVE classification (Chapter 3), this is the clearest case: GROUNDED mathematics producing SPECULATIVE physics with zero empirical social consequence. We keep it in the engine because it's correct and beautiful. We report it here because honesty requires admitting when beauty is useless.

13.3 What Dimension Teaches

The 3D and curvature results together make a point about the engine's cooperation mechanism: it is *topological*, not *geometric*.

Cooperation depends on connectivity (can agents reach each other?) not on metric (what shape is the space?). In 2D and 3D, agents that can reach resonant partners cooperate; agents that can't, don't. The specific geometry — flat vs curved, 2D vs 3D — changes how quickly agents find each other but not whether they cooperate once they do.

This explains why the gradient topology experiment (Finding 41) found cooperation robust to structural changes: fragmented perception changed the *number* of clusters (5.9 vs 1.2) but not the *fact* of clustering. Obstacles, moving wells, and short sight all changed geometry without changing topology — agents could still reach *some* partners, and some was enough.

The only topology change that kills cooperation is disconnecting the social graph entirely — making it impossible for any agent to reach any resonant partner. This requires either zero harmony range (Finding 33: range is CRITICAL) or zero resonance

regeneration (Finding 32: NO_REGEN kills 80%). Every other structural modification preserves the topology and therefore preserves cooperation.

Cooperation is a topological invariant of the engine. Geometry is decoration.

Next: Chapter 13 — What Simulations Can and Cannot Tell Us

Part III

The Unraveling

Chapter 14

Chapter 13: The Thread That Unraveled Everything

In which we ask the question we should have asked first, and discover that everything we believed was wrong.

14.1 The Question

Forty-five findings. Fifty-five experiments. A thesis that felt unassailable: cooperation is thermodynamically inevitable under energy scarcity with social range and mutual benefit.

Then a reviewer asked: “Did cooperation emerge because you found a real mechanism, or because you built the answer into the movement rule?”

We had never tested this. In every experiment from Chapter 5 through Chapter 12, agents followed the same FEP-style behavioral gradient — a handcrafted function that explicitly directs agents toward resonant partners. We had varied the environment in dozens of ways — scarcity, adversaries, topology, volatility, dimensionality — but we had never varied the *controller*. We had tested our thesis against every environmental challenge we could imagine while leaving the one variable that mattered completely untouched.

This is the most common error in simulation science: testing your hypothesis against perturbations while leaving the mechanism fixed. If the mechanism is designed to produce cooperation, then cooperation will appear in every environment — not because cooperation is inevitable, but because you hardcoded it into the agent’s behavior.

We built the experiment in an afternoon. Six controllers, identical thermodynamics:

1. **FEP_GRADIENT**: our original controller — seek partners, seek wells, flee danger, explore
2. **WELL_ONLY**: navigate to the nearest energy well, ignore other agents entirely
3. **GREEDY**: look in eight directions, move toward maximum energy gain
4. **PARTNER_ONLY**: navigate to the nearest resonant partner, ignore wells
5. **RANDOM**: uniform random direction each tick

6. STATIONARY: don't move at all

Same energy budgets. Same maintenance costs. Same resonance mechanics. Same wells. Same harmony profiles. Same 20 seeds. Only the movement rule changed.

14.2 The Result

Every controller produced cooperation.

WELL_ONLY agents — who never seek social contact, who navigate exclusively toward energy wells — cooperated 1.22 million times across 8,000 ticks. More than the FEP gradient's 1.08 million. They survived at 20 out of 20. The FEP gradient survived at 10.8.

STATIONARY agents — who do not move at all — cooperated 521,000 times. Twenty out of twenty survived.

RANDOM agents — walking in arbitrary directions with no goal — cooperated 474,000 times. Twenty out of twenty survived.

The cooperation we had spent twelve chapters analyzing, testing, and celebrating was not caused by our carefully designed behavioral gradient. It was caused by *being near energy wells at the same time as other agents*. Any controller that brings agents near wells — or that simply doesn't move them away — produces cooperation as a side effect of spatial proximity.

14.3 The FEP Gradient Was the Problem

The most disturbing column in the data was survival. Our FEP gradient — the sophisticated, multi-component, carefully weighted behavioral rule that we had built the entire book around — produced the *worst survival of any controller*.

10.8 out of 20. Worse than random walking. Worse than sitting still.

The mechanism was clear once we saw it. The FEP gradient has a social component that pulls agents toward resonant partners. This component diverts agents from energy wells — the actual source of survival — to pursue social contact. The social contact generates resonance energy, which partially compensates for the energy lost by being away from wells. But the compensation is incomplete. Agents following the FEP gradient spend ticks traveling between wells and partners, burning energy on movement that well-sitting agents conserve.

The gradient creates the problem that cooperation then solves. Without the social component, there is no problem — and no need for the solution.

14.4 What By-Product Mutualism Means

In behavioral ecology, *by-product mutualism* describes a situation where one organism's selfish actions incidentally benefit others nearby. A vulture circling a carcass attracts other vultures — not through communication or cooperation, but because the circling is visible and other vultures are watching. The first vulture doesn't intend to share; it intends to eat. The social structure (a group of vultures at a carcass) is a by-product of individual resource-seeking.

Our engine reproduces this dynamic with mathematical precision. Agents seek wells because wells provide energy. Wells are spatial points. Multiple agents seeking the same spatial points co-locate. Co-located agents passively resonate because resonance is proximity-based. The resonance provides bonus energy — real, measurable, incidental. No agent sought the resonance. No agent needed the resonance. The resonance was free.

This is by-product mutualism: selfish resource-seeking produces incidental social structure as a geometric consequence of shared spatial attractors.

Next: Chapter 14 — Seven Failed Rescues

Chapter 15

Chapter 14: Seven Failed Rescues

In which we try every environmental modification we can imagine to save our thesis, and each one confirms the same uncomfortable truth.

15.1 The Rescue Attempts

Finding 46 did not kill our thesis cleanly. It raised a possibility — that cooperation was by-product mutualism — but left room for the counterargument that the *environment* was too easy. Maybe cooperation becomes necessary under harsher conditions. Maybe the FEP gradient becomes indispensable when resources are scarce, volatile, or hidden.

We spent seven experiments testing this. Each was designed to find the environmental regime where the FEP gradient’s social component transitions from liability to necessity. Each failed.

15.2 Rescue 1: Make Wells Scarce (F48)

If cooperation is unnecessary because wells are too generous, reduce the generosity. We swept well capacity from 2,500 joules to 250 joules — a ten-fold reduction.

WELL_ONLY agents survived at 20 out of 20 at every capacity. Even at 250 joules — one-tenth of the original — purely selfish well-seeking was sufficient. The FEP gradient remained the worst controller at every scarcity level.

The wells were not too generous. A single agent standing at a well, even a tiny one, regenerates enough energy to survive indefinitely. The well regeneration rate (0.12 joules per tick) exceeds the per-agent share of maintenance costs. Scarcity doesn’t help because individual well access is always sufficient.

15.3 Rescue 2: Make Wells Volatile (F50)

If static wells are too predictable, make them cycle on and off. We tested four volatility regimes: stable (always on), seasonal (alternating 1,500-tick on/off cycles), random (50% chance of flipping every 500 ticks), and migrating (wells physically relocate every 1,500 ticks).

Under stable, seasonal, and random volatility, WELL_ONLY agents survived at 20 out of 20. Agents simply waited at well locations during off-periods, surviving on initial energy reserves and occasional resonance from nearby waiting agents.

Migrating resources produced the only marginal result in the entire rescue series. When wells physically moved, WELL_ONLY survival dropped to 9.3 and FEP survival matched at 9.9 — a small, non-significant convergence ($d = 0.35$). For the first time, the social gradient was not a clear liability. But it was not a clear asset either. Both controllers suffered equally under migration because both had to search for relocated wells.

15.4 Rescue 3: Make Wells Hidden (F49)

If agents can find wells too easily, make the wells invisible. We tested four visibility levels: full perception (200 units), short-range (30 units), completely hidden (0 units — agents must stumble onto wells), and a social information condition (wells invisible but high-energy agents attract the FEP gradient).

WELL_ONLY agents survived at 20 out of 20 even when wells were completely hidden. The arena was 100 by 100 units. The well interaction radius was 35 units. A random walker in a 100-unit arena encounters a 35-unit-radius target frequently enough to survive. The arena was simply too small for hidden resources to matter.

The FEP gradient performed worse than WELL_ONLY under every visibility condition, including hidden (16.8 vs 20.0, $d = 1.39$). Even when wells were invisible, the social component hurt more than it helped.

15.5 Rescue 4: Require Threshold Co-Location (F51)

If cooperation fails because individual well access is too easy, make wells require multiple agents. We implemented a threshold mechanism: wells only dispense energy when two, three, or four agents are simultaneously present within the well radius.

WELL_ONLY agents survived at 20 out of 20 at every threshold, including the strictest (four agents required). The mechanism: twenty agents seeking three wells naturally produces six to seven agents per well. The threshold was always satisfied as a side effect of shared attraction to the same spatial points.

The FEP gradient performed worst again, because its social component scattered agents across the arena rather than concentrating them at wells. At threshold four, FEP managed 10.8 out of 20 while WELL_ONLY achieved a perfect 20.

15.6 Rescue 5: Require Compatible Partners (F52)

If the threshold test failed because any agents co-locating is sufficient, require *compatible* partners. We implemented resonance-gated wells: a well only provides energy to an agent if that agent has a partner at the same well whose harmony resonance exceeds 0.5 or 0.7.

The GREEDY controller dominated at 16.5 to 16.9 across all resonance gates. Its eight-direction lookahead naturally optimized for positions with both well access and resonant partners. FEP achieved 10.9 at the strictest gate (0.7), marginally better than WELL_ONLY's 9.3, but the difference was not significant ($d = -0.20$, $p = 0.57$).

The problem: with twenty agents drawn from six harmony profiles, the probability of at least one compatible pair among six to seven agents at a well is high. Dense well-clustering produces compatible pairs by chance, without any agent seeking compatibility.

15.7 Rescue 6: Remove Wells Entirely (F53)

The nuclear option. If wells make individual survival too easy, remove them. Make resonance — social energy from compatible proximity — the only energy source. Surely, in this regime, social seeking would become indispensable.

It did not. STATIONARY agents — who never move — survived at 19.5 out of 20. FEP agents survived at 11.4. The FEP gradient was *significantly worse* than doing nothing ($d = 1.80$, $p < 0.0001$).

The mechanism: in an 80-by-80 arena with 20 agents, initial random placement produces enough proximity for passive resonance encounters. Agents that sit still conserve the energy that moving agents waste on locomotion. The initial energy budget (200 joules at 0.15 per tick = 1,333 ticks of buffer) provides ample time for passive resonance to accumulate.

The GREEDY controller survived at 19.9 by optimizing direction for resonance proximity — but STATIONARY was nearly as good, confirming that in this arena, optimal behavior is *inaction*.

15.8 Rescue 7: The FEP Gradient as Informational Drive

Across all seven rescues, a pattern emerged. The FEP gradient's social component is not a cooperation mechanism. It is an *informational drive* — a compulsion to seek social contact that competes with basic metabolic efficiency. In environments where metabolic efficiency determines survival, the informational drive is a liability.

The drive becomes approximately neutral (not beneficial — just not harmful) only under one condition: migrating resources (F50, $d = 0.35$). When resource locations change, the social gradient provides weak information about where other agents are moving, which correlates weakly with where resources have appeared. This is the producer-scrouter dynamic from behavioral ecology: the social gradient enables information scrounging from agents who have found resources.

But even in this best case, the social gradient does not *outperform* well-seeking. It merely converges with it. The FEP gradient's social component is, at its very best, neutral.

15.9 What the Rescues Taught Us

Seven experiments. Seven failures. Each confirmed the same mechanism from a different angle:

1. Individual resource access is always sufficient (F48)
2. Temporal gaps don't require cooperation; they require waiting (F50)

3. Spatial uncertainty doesn't require social information; it requires random exploration (F49)
4. Synchronized access is free when everyone wants the same thing (F51)
5. Compatible pairing is free when populations are dense (F52)
6. Even without resources, passive proximity suffices (F53)
7. The social drive hurts more than it helps in every regime

The thesis cannot be rescued because the thesis was wrong. Cooperation in this engine is not a mechanism; it is a geometry.

Next: Chapter 15 — Resource Geometry: The True Determinant of Social Structure

Chapter 16

Chapter 15: Resource Geometry

In which we state the revised thesis plainly and ask what it means.

16.1 The Thesis, Revised

After 53 findings and 63 experiments, the thesis is:

Within the tested regime — small arenas, locatable wells, non-monopolizable access, passive proximity benefit — the spatial patterns we measured (clustering, incidental energy transfer, co-location, survival distributions) are determined by the geometry of resource attractors, not by the movement algorithms of agents.

Six distinct controllers — including agents that do not move at all — produce all the spatial-statistical phenomena we originally attributed to a cooperation mechanism. Clustering, resonance events, co-location patterns, and survival distributions are geometric consequences of shared spatial attractors. We use the term “social structure” for readability but acknowledge it labels a spatial-statistical phenomenon, not a social one in the human sense.

This is by-product mutualism: selfish action produces incidental mutual benefit when the environment channels individuals toward the same locations.

16.2 What This Explains

Every finding from Chapters 5 through 12 is reinterpreted under resource geometry:

Clustering (F1-4) emerges because all agents seek the same two or three wells, not because the FEP gradient creates social structure. Any controller that navigates to wells produces identical clustering.

Phase transitions (F20) occur because maintenance pressure depletes energy faster than wells can regenerate it. The transition is between well-access regimes, not cooperation regimes. Below the critical pressure, well regen exceeds maintenance. Above it, the deficit grows faster than resonance can compensate.

Social bonds (F19) improved survival by 70% because distance constraints held agents near wells, not near each other. The bond prevented drift away from resource access points.

Adversarial dynamics (F14-18) reflected competition for well access, not social conflict. Adversaries died faster because their inverted harmony profiles meant they gained no incidental resonance — not because defection was punished, but because they received less incidental benefit from co-location.

The equal opportunity paradox (F26) was a race for well access, not a cooperation failure. Agents starting at the same point all competed for the nearest well, and small stochastic advantages compounded. This is resource competition, not social inequality.

The Bowling Alone result (F43) showed that voluntarily refusing to co-locate with other agents reduces incidental resonance. Under the revised thesis, “bowling alone” means voluntarily walking away from resource points where others are already gathering — giving up the free benefit of by-product mutualism, not refusing to cooperate.

16.3 What This Does Not Explain

The engine cannot speak to environments where:

- Resources are non-spatial (information, social capital, reputation)
- Resources require genuine coordination to extract (pack hunting, collective construction)
- Agents have conflicting resource preferences (not all wanting the same wells)
- The arena is large enough that random exploration fails
- Agents can monopolize resources (excluding others from wells)

These are the conditions under which real cooperation — not by-product mutualism — becomes necessary. Our engine’s design excludes all of them, which is why it found only by-product mutualism. This is not a flaw in the engine; it is a precise statement of the engine’s scope.

16.4 The Constructive Proof, Revised

Chapter 13 of the original manuscript argued that the work had value as a constructive proof of sufficiency: “These conditions are sufficient to produce cooperation.”

The revised constructive proof is more interesting: “These conditions are sufficient to produce *the appearance* of cooperation without any cooperation mechanism.” The social structure is real. The clustering is real. The survival patterns are real. The cooperation events are real. But they are all by-products of individual resource-seeking, not evidence of a cooperation mechanism.

This is a contribution to simulation science: it demonstrates how easily emergent behavior can be misattributed when the researcher tests the environment but not the controller.

16.5 The Lesson for Simulation Science

We made a specific, correctable error. We varied everything except the one thing that mattered. We changed the environment 45 times, added adversaries, volatility, topology, dimensionality, learning, memory, communication — and concluded that our findings were robust. They were robust to everything except the one test we did not run: replacing the controller.

This is not unique to our project. Any agent-based model that uses a fixed behavioral rule and varies only environmental parameters is vulnerable to the same error. The behavior may be a property of the rule, not the environment. The only way to know is to replace the rule and see if the behavior survives.

Our by-product mutualism result survived controller replacement — cooperation appeared under all six controllers. But its *mechanism* did not survive. The mechanism we attributed to the FEP gradient was actually resource geometry. The falsification was productive: it produced a cleaner, more parsimonious, and more general explanation.

We recommend that every simulation paper include a controller replacement study as standard methodology. If your finding survives controller replacement, it is environment-driven. If it does not, it is controller-driven. Both are valid findings, but they mean different things.

16.6 The Closing

We began this book expecting to demonstrate that cooperation is thermodynamically inevitable. We demonstrated instead that social structure is geometrically inevitable — a weaker but more honest claim, and ultimately a more interesting one.

The geometry of resources determines the geometry of societies. Agents do not cooperate because they must. They co-locate because the wells are where the energy is. Everything else — the clustering, the apparent coordination, the social structure that looks so much like community — is free.

Physics does not cooperate. It clusters at the wells. And in that clustering, incidentally, it produces everything that looks like a society.

Next: Chapter 16 — The Freedom to Walk Away

Chapter 17

Chapter 16: The Freedom to Walk Away

In which the one finding that survives the revision tells us something about choice.

17.1 What Survived

Fifty-three findings. Seven that falsified the original thesis. Sixteen chapters rewritten. One finding survives — not unchanged in interpretation, but unchanged as a pattern.

Finding 43: when agents gain the freedom to sometimes ignore resource co-location — to walk away from wells where others are gathered — the incidental benefits of by-product mutualism collapse. Cooperation events drop 53%. Isolation increases. Inequality rises. And the worst outcome belongs to the half-committed: agents who sometimes co-locate and sometimes don't, paying the movement cost of both strategies while reliably accessing neither.

Under the original thesis, this was about cooperation and commitment. Under the revised thesis, it is about something different: **the freedom to leave a resource point**. The pattern is the same four Putnam trends. The mechanism has changed entirely — from “refusing to cooperate” to “refusing to stay at the well.”

17.2 By-Product Mutualism Requires Co-Location

By-product mutualism is automatic. It requires no intent, no strategy, no social drive. But it does require one thing: agents must be at the same place at the same time.

In our engine, agents co-locate at wells because wells are spatial attractors. Every controller that seeks wells produces co-location. Co-location produces passive resonance. Passive resonance provides incidental energy. The entire mechanism is automatic — as long as agents stay near wells.

The willingness parameter (F43) breaks the one condition that by-product mutualism requires. An agent with willingness 0.6 sometimes follows the gradient to wells and sometimes ignores it. On ignoring ticks, the agent wanders away from the well — away from other agents — and the passive resonance stops. The agent loses the incidental benefit not because the benefit was taken away, but because the agent chose to leave.

17.3 The U-Curve Is About Resource Access, Not Commitment

The U-shaped survival curve (full cooperation and full selfishness both outperform partial commitment) has a cleaner explanation under the revised thesis.

At willingness 1.0, agents consistently co-locate at wells. They receive both well energy and incidental resonance. Both energy streams are reliable because the agent's behavior is consistent.

At willingness 0.0, agents ignore all external signals and stay put or wander minimally. They don't waste energy on directed movement. In the resonance-only experiment (F53), stationary agents outperformed every active controller. The couch potato strategy works because movement costs energy and the arena is small enough that passive encounters provide sufficient resonance.

At willingness 0.6, agents oscillate between seeking wells and ignoring the gradient. They spend ticks moving toward wells, then ticks wandering away. They arrive at wells intermittently, accessing both energy streams unreliably. The movement cost is paid on every tick regardless. The result is worse than either consistent strategy.

This is not about social commitment. It is about behavioral consistency in resource access. An agent that reliably goes to wells thrives. An agent that reliably stays still thrives. An agent that inconsistently does both wastes energy on movement without reliable access to either energy stream.

17.4 The Bowling Alone Reinterpretation

Putnam's four trends — declining civic participation, increasing social isolation, persistent economic output, and rising inequality — all appeared at low willingness levels. Under the original thesis, this was about the failure of cooperative intent. Under the revised thesis, it is about the failure of consistent resource access.

Declining participation: agents access wells less frequently, reducing the incidental cooperation that wells produce as a side effect of co-location.

Increasing isolation: agents wander between wells and open space, spending more time alone.

Persistent economic output: individual survival persists because the wells still exist and agents still sometimes reach them. The economy (survival rate) doesn't collapse because individual resource access is still possible — just less efficient.

Rising inequality: inconsistent well access creates variance. Some agents happen to find wells on their wandering ticks; others don't. The stochastic nature of inconsistent behavior amplifies into distributional inequality.

The pattern is identical to Putnam's observations. The mechanism is different: not the decline of social capital, but the decline of consistent resource-access behavior.

17.5 What This Means

The revised thesis makes a specific prediction about the relationship between resources and social structure:

Social structure persists exactly as long as agents co-locate at shared resources. It collapses exactly when agents gain the freedom — and exercise the choice — to leave.

This is not a metaphor. It is a mathematical relationship between willingness to co-locate and the four measurable social outcomes (participation, isolation, survival, inequality). The relationship is monotonic between willingness 1.0 and 0.2, with the U-curve at 0.6 creating a local minimum where inconsistency is maximally destructive.

In the real world, this maps not to the decline of cooperative intent but to the decline of physical co-location. When people leave the factory floor, the church, the bowling league, the neighborhood bar — when they choose screens over proximity — they walk away from the spatial conditions that produce by-product mutualism. The social benefits they lose were never chosen in the first place. They were free. And they disappear the moment the conditions that produced them are no longer met.

17.6 The Last Word

We expected to write a book about cooperation. We wrote a book about geometry instead.

Every social structure we observed — the clustering, the cooperation, the phase transitions, the group formation, the solidarity, the inequality — was a geometric consequence of agents co-locating at shared resources. No social mechanism was required. No integration metric mattered. No cognitive architecture contributed. Physics does not cooperate. It clusters at the wells.

But one thing disrupts the clustering: choice. Give agents the freedom to walk away from the wells, and the geometry breaks. The by-product mutualism that produced all the appearance of society dissolves — not because cooperation failed, but because co-location ceased.

The thermodynamics of togetherness, it turns out, requires only one thing: being in the same place. And the tragedy of walking away is that you lose benefits you never knew you were receiving.

17.7 Appendix Note

All 63 experiments can be reproduced:

```
cargo run --example [name] --release
```

Seeds are deterministic. Results are bitwise reproducible. Statistical methods, effect sizes, and corrections are computed inline.

If the prose and the data disagree, trust the data. We did.

Appendix B: Old Interpretation vs Revised Interpretation

Finding	Original Interpretation (Part II)	Falsifying Evidence
F1-4: Clustering	FEP gradient drives social clustering	F46: all controllers cluster
F19: Social bonds +70%	Friendship springs sustain cooperation	F46: WELL_ONLY survive
F20: Phase transition	Cooperation/collapse boundary	F47: wells alone sustain
F24: Abundance curse	Cooperation pressure removed	Unchanged
F25: Altruism fails	Charity without structure fails	Unchanged (strengthened)
F26: Equal starts → inequality	Cooperation failure	F46: resource competition
F14-18: Adversaries	Social conflict; solidarity response	F47: resonance is incidental
F29: Info irrelevant	FEP already optimal	F46: all controllers work
F30: 3D cooperation	Cooperation generalizes to 3D	Unchanged
F31: Weights optimal	Cooperation already tuned	F46: weights don't matter
F34: Memory useless	FEP encodes all info	F46: gradient unnecessary
F36: Curvature irrelevant	Math correct, socially inert	Unchanged
F38: Can't kill cooperation	Thermodynamic inevitability	F46-48: any controller survive
F40: Communication zero effect	FEP already optimal	F46: no gradient needed
F43: Bowling Alone	Voluntary cooperation refusal	Reinterpreted
F44: U-curve mechanism	Partial commitment worst	Reinterpreted

Afterword: What This Changes

A personal note from the author about what 53 findings mean for the systems we actually build.

A reviewer asked: “Given that social structure is dictated by resource geometry rather than internal algorithms, how does this change how you design real-world infrastructure?”

It changes everything.

I spent two years building Mycelix — a decentralized civic operating system with 16 application clusters, 134 Holochain zones, and a consciousness-gating system that adjusts permissions based on an integration metric. The implicit assumption behind all of it was that if you build the right social mechanisms — the right voting protocols, the right reputation systems, the right harmony profiles — people will cooperate.

Finding 46 says that assumption is wrong. Or at least incomplete.

People cooperate when the resource geometry forces co-location. They cooperate at the water cooler, at the school gate, at the neighborhood bar — not because they chose to, but because they’re there. The social structures that emerge from these encounters are by-product mutualism: real benefits from incidental proximity.

The implication for system design is not to build better social algorithms. It’s to build better wells.

A community platform that optimizes for engagement (keeping people on the app) is optimizing the wrong variable. The simulation showed that movement toward social contact is a metabolic cost — it pulls agents away from resources. A platform that optimizes for co-location at shared resources (community spaces, shared tools, collective projects) produces cooperation as a free side effect.

Concretely:

Design the wells, not the gradient. Instead of recommending connections (a social gradient), create shared resources that people must physically or digitally co-locate to access. A community workshop. A shared compute cluster. A collective garden. The cooperation will happen incidentally.

Make the wells non-monopolizable. Finding 51 showed that threshold access (requiring multiple agents) doesn’t help when everyone converges on the same well anyway. But monopolizable resources — where one agent can exclude others — would create genuine scarcity that might require genuine coordination. The real design challenge is resources that are shared but contestable.

Don’t build social drives. The FEP gradient’s social component was a metabolic liability in every regime we tested. Translating this to platform design: social features that nudge people toward interaction (notifications, recommendations, “people you may know”) are the digital equivalent of the social gradient. They create movement that costs energy without necessarily producing co-location at resources.

Protect co-location from voluntary departure. Finding 43 showed that the freedom to walk away collapses by-product mutualism. In platform terms: make it easy to stay at the well (low friction, high accessibility) and expensive to leave (not through lock-in, but through genuine value that accrues from presence). The tragedy is not that people leave — it's that they don't know what they're losing.

This is a humbling conclusion for someone who built a consciousness-gating system. The simulation says consciousness metrics don't matter (Finding 3). Social drives don't help (Finding 46). Communication adds nothing (Finding 40). Memory is irrelevant (Finding 34). The only thing that matters is whether agents are near shared resources.

The Eight Harmonies that structure Mycelix — Stillness, Play, Craft, Justice, Curiosity, Celebration, Kinship, Stewardship — may be beautiful as a philosophical framework. But the simulation suggests they function as harmony profiles that enable passive resonance, not as active drivers of cooperation. Their value is incidental, not causal.

I don't know yet what to do with this. It may mean redesigning Mycelix around shared resources rather than social mechanisms. It may mean the consciousness-gating system is unnecessary overhead — a metabolic tax on participation, like the FEP gradient. It may mean the entire paradigm of “consciousness-first technology serving all beings” needs to become “resource-first technology that incidentally produces community.”

That is uncomfortable. It is also what the data says. And the last line of this book applies to the author as much as the reader:

If the prose and the data disagree, trust the data. We did.

Reproduction

All 63 experiments reported in this book can be reproduced:

```
cargo run --example [experiment_name] --release
```

Seeds are deterministic. Results are bitwise reproducible on the same platform (verified by lock-in test). Statistical methods, effect sizes, and Holm–Bonferroni corrections are computed inline.

63 experiment files. 53 findings. 95 library tests + 2 integration tests. AGPL-3.0.

If the prose and the data disagree, trust the data. We did.